

TensorFlow Tutorial

Deep Learning School 2016

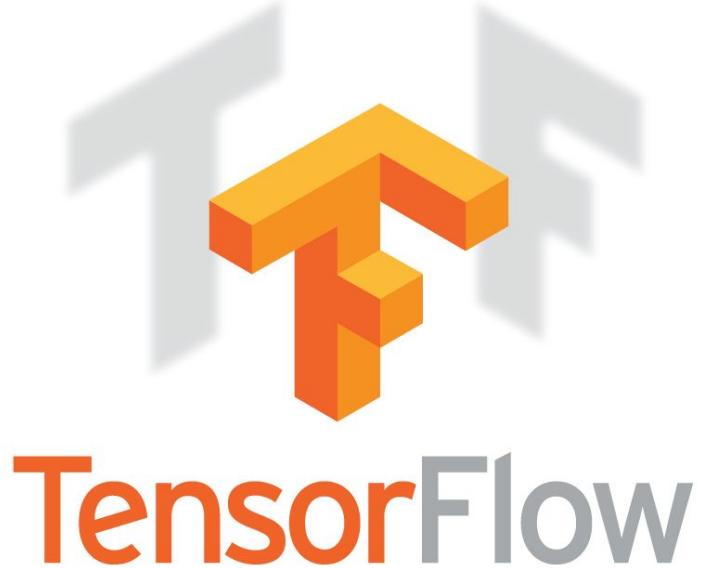


Sherry Q. Moore

Have you installed TensorFlow?

<https://github.com/sherrym/tf-tutorial/>

```
# cd tf-tutorial  
# git pull  
# source ~/tensorflow/bin/activate  
# jupyter notebook
```

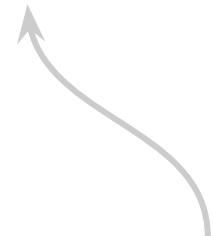


- **Open source** Machine Learning library
- Especially useful for **Deep Learning**
- For research **and** production
- **Apache 2.0** license

A multidimensional array.



TensorFlow

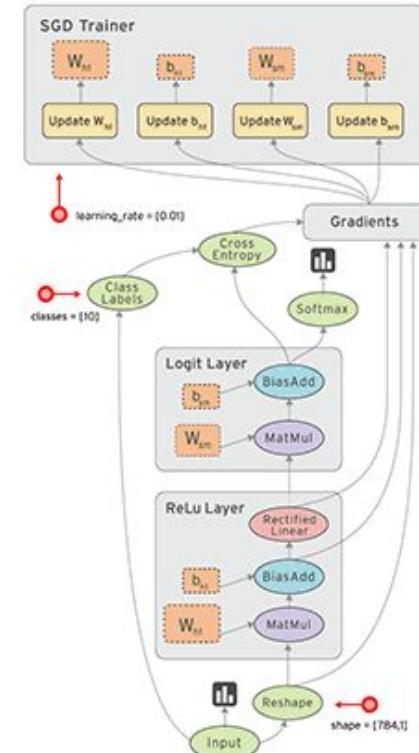


A graph of operations.

Data Flow Graphs

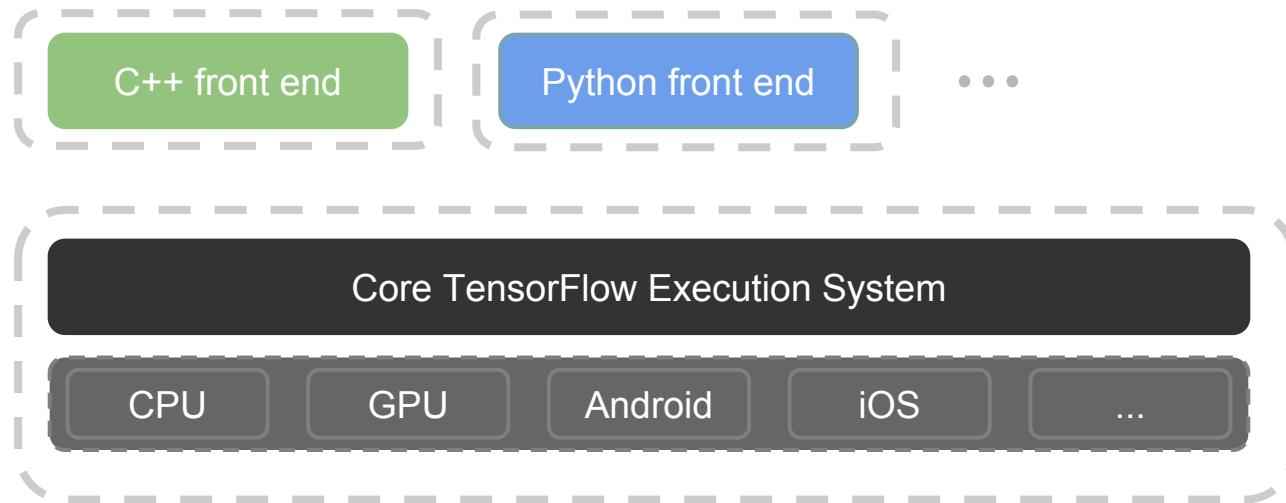
Computation is defined as a directed acyclic graph (DAG) to optimize an objective function

- Graph is defined in high-level language (Python)
- Graph is compiled and optimized
- Graph is executed (in parts or fully) on available low level devices (CPU, GPU)
- Data (tensors) flow through the graph
- TensorFlow can compute gradients automatically



Architecture

- Core in C++
- Different front ends
 - Python and C++ today, community may add more



Portable & Scalable



Your laptop



Datacenters



Android



iOS



Raspberry
Pi

Models on TensorFlow

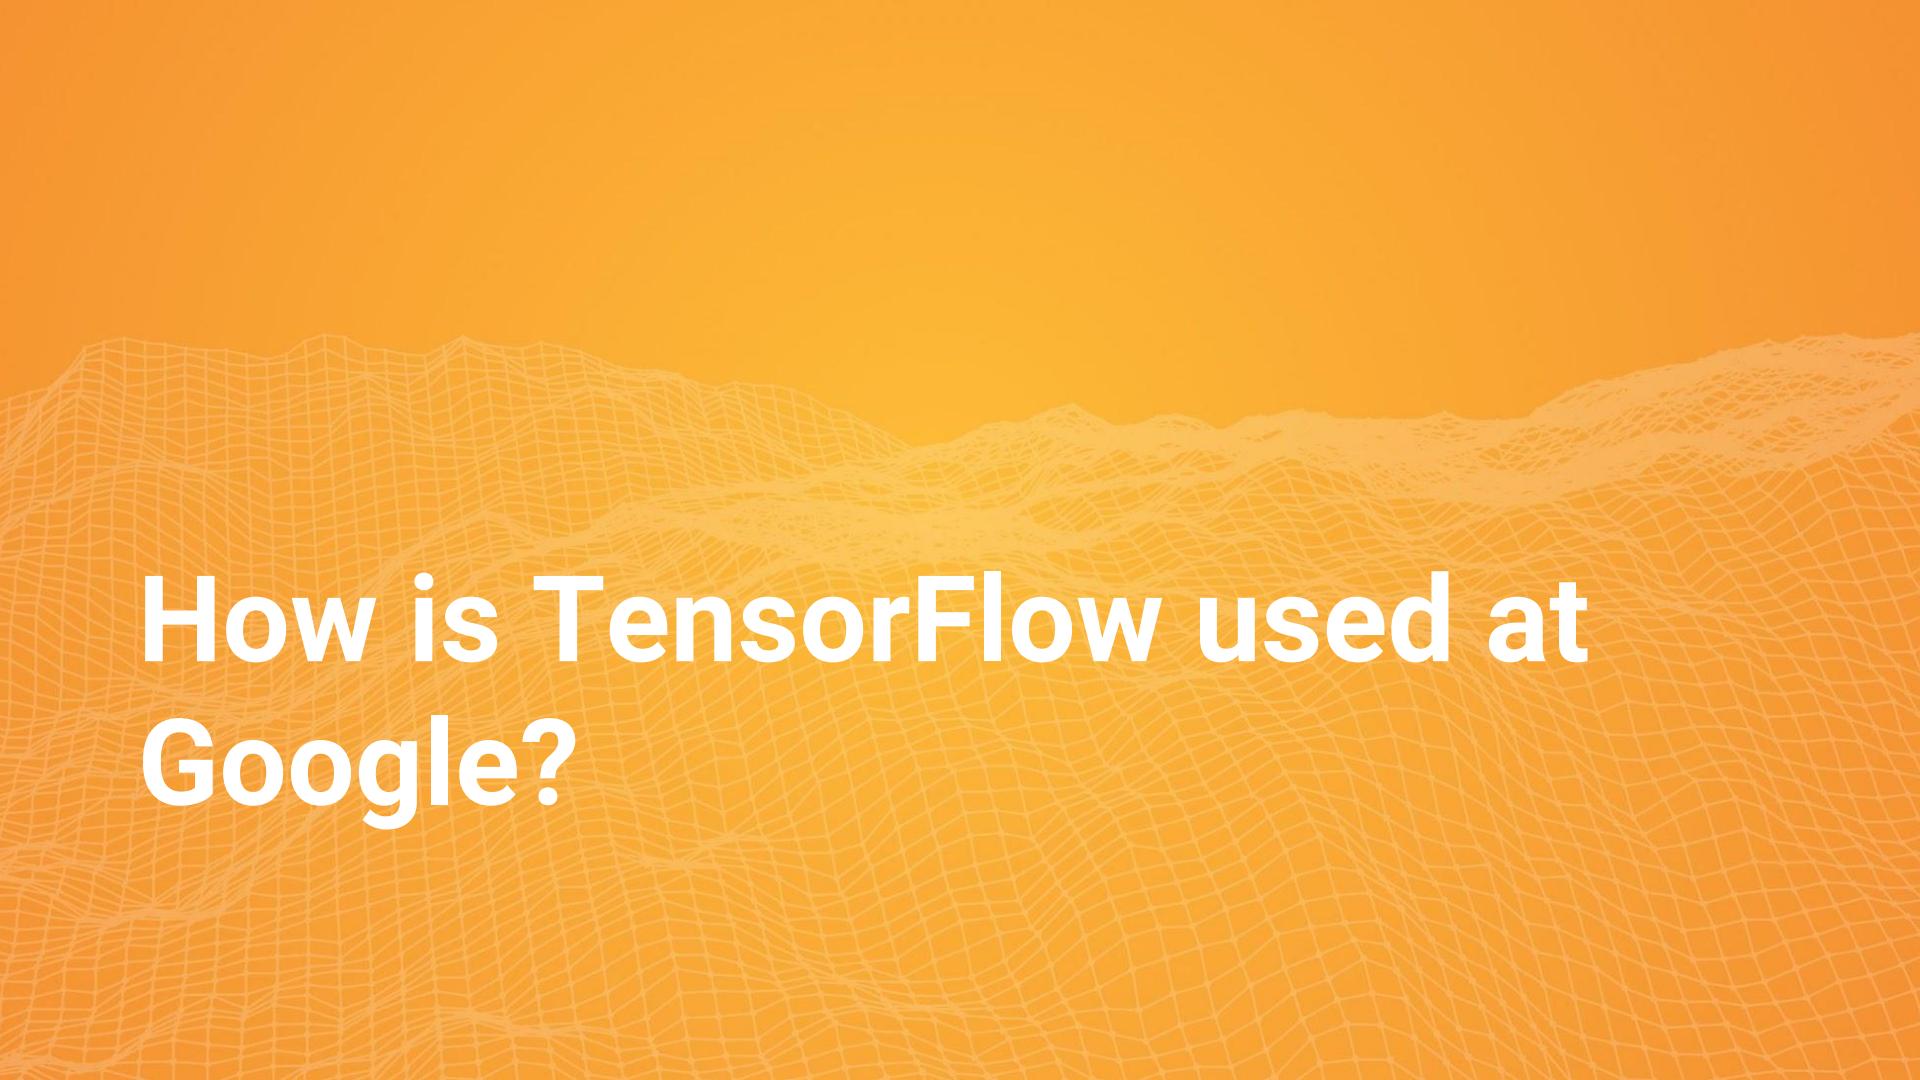
<https://github.com/tensorflow/models>

- Inception <https://github.com/tensorflow/models/tree/master/inception>
- Show and Tell: A Neural Image Caption Generator
<https://github.com/tensorflow/models/tree/master/im2txt>
- Language Model on One Billion Word
https://github.com/tensorflow/models/tree/master/lm_1b
- SyntaxNet: Neural Models of Syntax
<https://github.com/tensorflow/models/tree/master/syntaxnet>
- Resnet on Cifar10 and Cifar100
<https://github.com/tensorflow/models/tree/master/resnet>
- Sequence-to-Sequence with Attention Model for Text Summarization
<https://github.com/tensorflow/models/tree/master/textsum>

High Level Libraries

<https://github.com/tensorflow/models>

- TensorFlow-Slim image classification library
<https://github.com/tensorflow/models/tree/master.slim>
- Image Compression with Neural Networks
<https://github.com/tensorflow/models/tree/master/compression>
- Autoencoders <https://github.com/tensorflow/models/tree/master/autoencoder>
- Swivel <https://github.com/tensorflow/models/blob/master/swivel>
- Spatial Transformer Network
<https://github.com/tensorflow/models/tree/master/transformer>

The background of the slide features a solid orange gradient at the top transitioning into a white wireframe mesh that resembles undulating hills or waves.

How is TensorFlow used at Google?

Recognizing images with Inception



spotted salamander



fire salamander

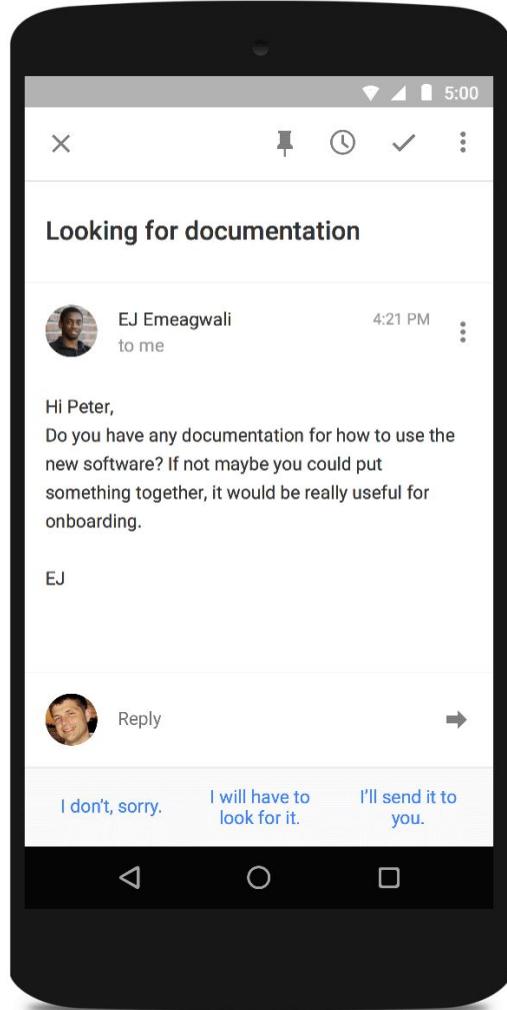


retriever



“Ok Google... Dog videos”





Smart reply in Inbox by Gmail

10%

of all responses
sent on mobile

AlphaGo

BBC News Sport Weather iPlayer TV Radio More Search

Find local news

Home UK World Business Politics Tech Science Health Education Entertainment & Arts More

Technology

Google achieves AI 'breakthrough' at Go

An artificial intelligence program developed by Google beats Europe's top player at the ancient Chinese game of Go, about a decade earlier than expected.

© 27 January 2016 | Technology

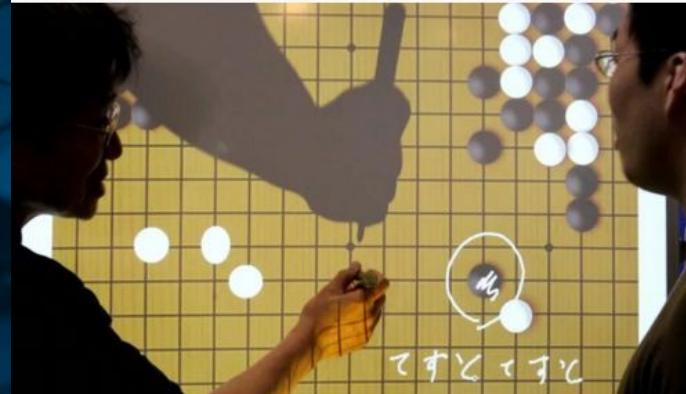
- How did they do it?
- What is the game Go?

Facebook trains AI to beat humans at Go



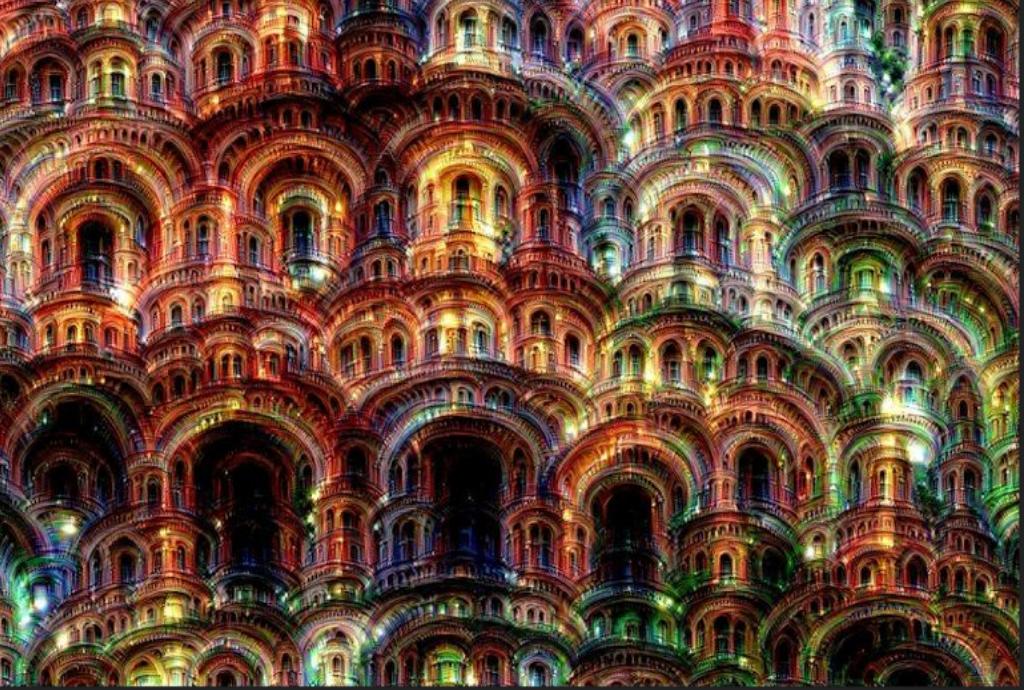
Google's AI just cracked the game that supposedly no computer could beat

By Mike Murphy | January 27, 2016



Going up. (Reuters/Kiyoshi Ota)

Computers have slowly started to encroach on activities we previously believed only the brilliantly sophisticated human brain could handle. IBM's Deep Blue supercomputer beat Grand Master Garry Kasparov at chess in 1997, and in 2011 IBM's Watson beat former human winners at the quiz game *Jeopardy*. But the ancient board game Go has long been one of the major goals of artificial intelligence research. It's understood to be one of the most difficult games for computers to handle due to the sheer number of possible moves a player can make at any given point. Until now, that is.



Lab#1 Linear Regression Model

https://github.com/sherrym/tf-tutorial/blob/master/1_linear_regression_model.ipynb

What are we trying to do?

Mystery equation: $y = 0.1 * x + 0.3$

Model: $y = w * x + b$

Objective: Given enough (x, y) value samples,
figure out the value of w and b .

1_linear_regression_model.ipynb

1. Import libraries (cell 1.1)
2. Create input data (cell 1.2)
3. Build inference graph (cell 1.3)
 - a. Create Variables
 - b. Create Operations that produce logistic outputs.
4. Build training graph (cell 1.4)
 - a. Loss
 - b. Optimizer
5. Create session and run initialization (cell 1.5)
 - a. Initialize all variables
 - b. Run train_op -- the operation that minimizes loss
6. Perform training (cell 1.6)

https://github.com/tensorflow/tensorflow/blob/master/tensorflow/g3doc/api_docs/python/index.md

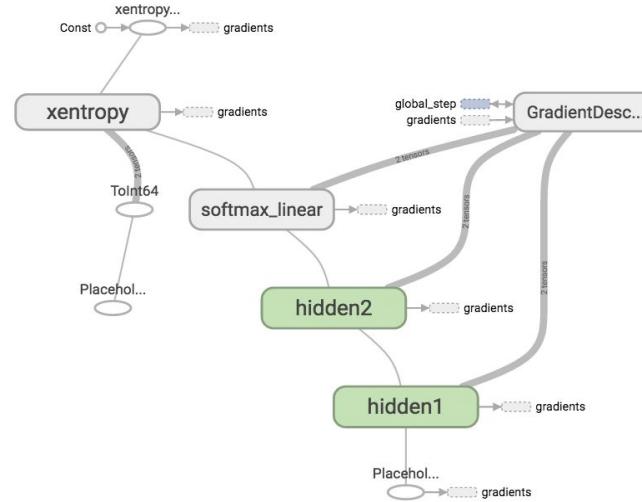
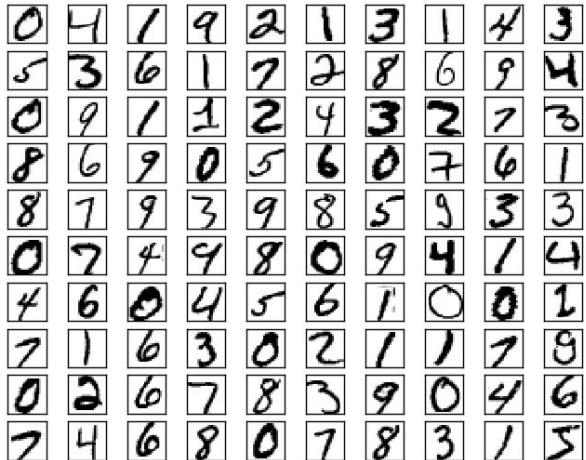
Exercises to run

1. Run all the cells
2. Uncomment 1.5 and see the graph
3. Uncomment 1.8 and compare values
4. Try to train a different linear function
 - If you can't visualize the values, try changing pylab.ylim()
5. Try a different initialization function.
6. Try a different optimizer.

Lab#2 MNIST

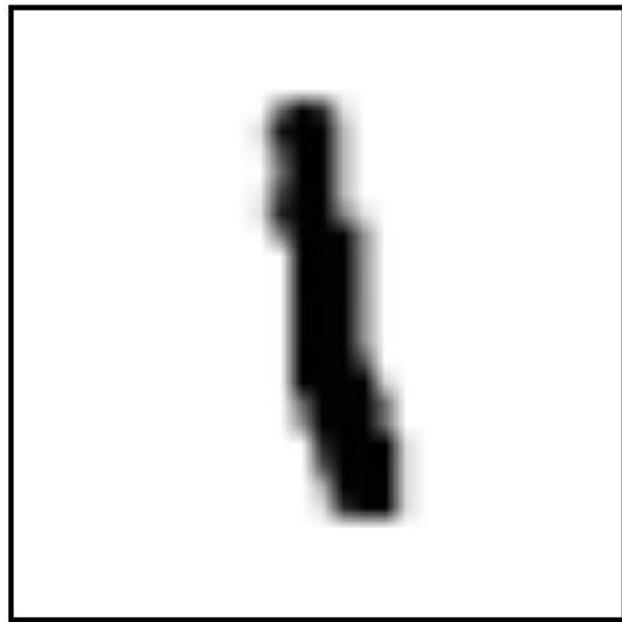
https://github.com/sherrym/tf-tutorial/blob/master/2_mnist.ipynb

What are we trying to do?



Objective: Given enough images and labels, figure out the weights and biases so the model can correctly identify digits.

What we see



2

What the computer “sees”

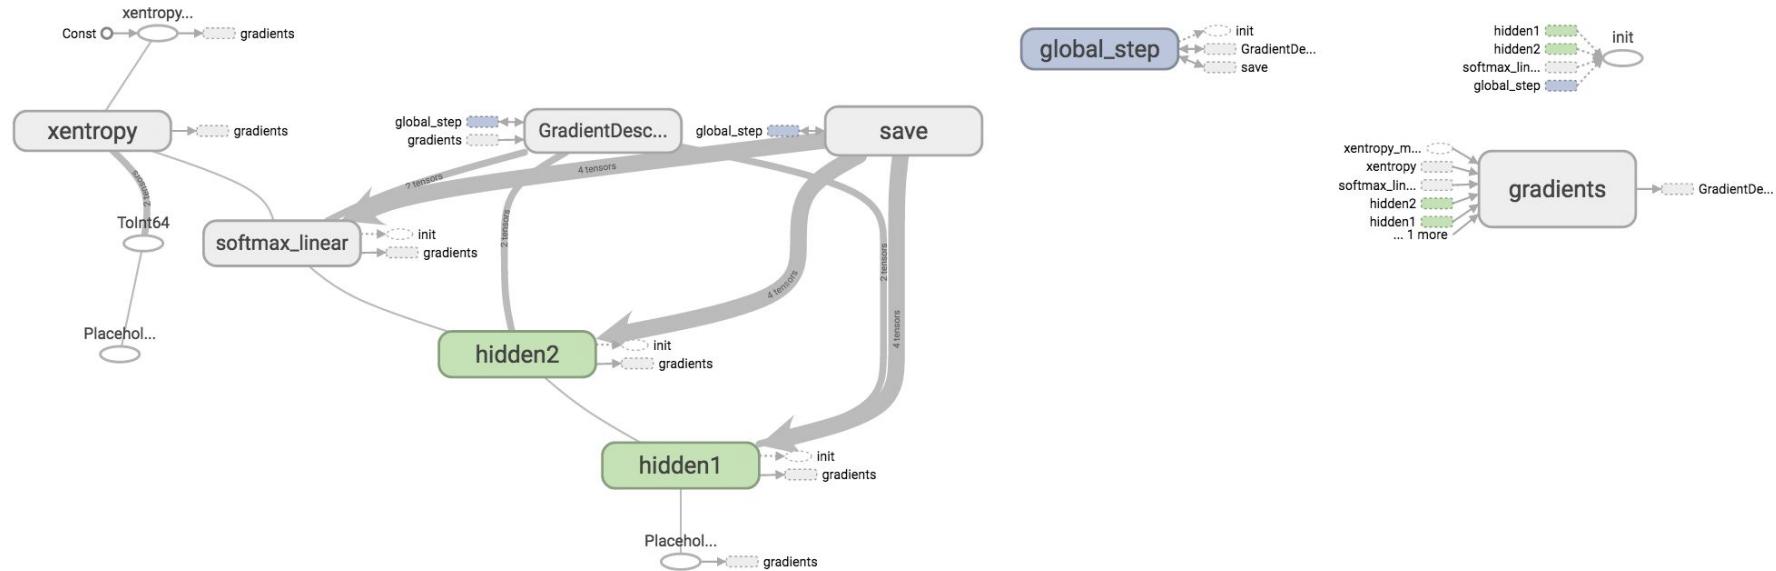
2_mnist.ipynb

1. Import libraries (cell 2.1)
2. Define a set of constants for convenience (cell 2.2)
3. Get input data (cell 2.3)
4. Build inference graph (cell 2.4)
5. Build training graph (cell 2.5)
6. Build the complete graph for training, and saving checkpoints (cell 2.6)
7. Run training for MAX_STEPS and save checkpoint at the end (cell 2.7)
8. Run evaluation based on the saved checkpoint (cell 2.8)

https://github.com/tensorflow/tensorflow/blob/master/tensorflow/g3doc/api_docs/python/index.md

Exercises to run

1. Run all the cells
2. Uncomment the graph writing to see what the graph looks like.
3. Print out loss value every 100 steps. What do you notice?
4. Save checkpoints every 500 steps.
5. Run evaluation with different checkpoints. What do you notice?
6. Run evaluation on the complete validation set.
7. Build evaluation graph from scratch instead of importing from meta graph.



Exercise: Evaluation on the whole set

```
# Build graph to count correct answers.
def mnist_evaluation(logits, labels):
    """Evaluate the quality of the logits at predicting the label.

    Args:
        logits: Logits tensor, float - [BATCH_SIZE, NUM_CLASSES].
        labels: Labels tensor, int32 - [BATCH_SIZE], with values in the
            range [0, NUM_CLASSES].
    Returns:
        A scalar int32 tensor with the number of examples (out of batch_size)
        that were predicted correctly.
    """
    # For a classifier model, we can use the in_top_k Op.
    # It returns a bool tensor with shape [batch_size] that is true for
    # the examples where the label is in the top k (here k=1)
    # of all logits for that example.
    correct = tf.nn.in_top_k(logits, labels, 1)
    # Return the number of true entries.
    return tf.reduce_sum(tf.cast(correct, tf.int32))
```

```
# Evaluation on the whole evaluation set.  
with tf.Session(graph=tf.Graph()) as sess:  
    saver = tf.train.import_meta_graph(  
        os.path.join(TRAIN_DIR, "checkpoint-1999.meta"))  
    saver.restore(  
        sess, os.path.join(TRAIN_DIR, "checkpoint-1999"))  
  
    print('Validation Data Eval:')  
    # And run one epoch of eval.  
    true_count = 0 # Counts the number of correct predictions.  
    steps_per_epoch = data_sets.validation.num_examples  
    num_examples = steps_per_epoch * BATCH_SIZE  
    # Retrieve the Ops we 'remembered'.  
    logits = tf.get_collection("logits")[0]  
    images_placeholder = tf.get_collection("images")[0]  
    labels_placeholder = tf.get_collection("labels")[0]  
  
    # Add the Op to compare the logits to the labels during evaluation.  
    eval_op = mnist_evaluation(logits, labels_placeholder)
```

```
for step in xrange(steps_per_epoch):  
    images_feed, labels_feed =  
        data_sets.validation.next_batch(BATCH_SIZE)  
    true_count += sess.run(eval_op,  
                          feed_dict={images_placeholder:  
                                      images_feed,  
                                      labels_placeholder: labels_feed})  
precision = true_count * 1.0 / num_examples  
print(' Num examples: %d Num correct: %d Precision @ 1:  
%0.04f %'  
      (num_examples, true_count, precision))
```

What's Next

tensorflow.org

github.com/tensorflow

Want to learn more?

Udacity class on Deep Learning, goo.gl/iHssl

Guides, codelabs, videos

MNIST for Beginners, goo.gl/tx8R2b

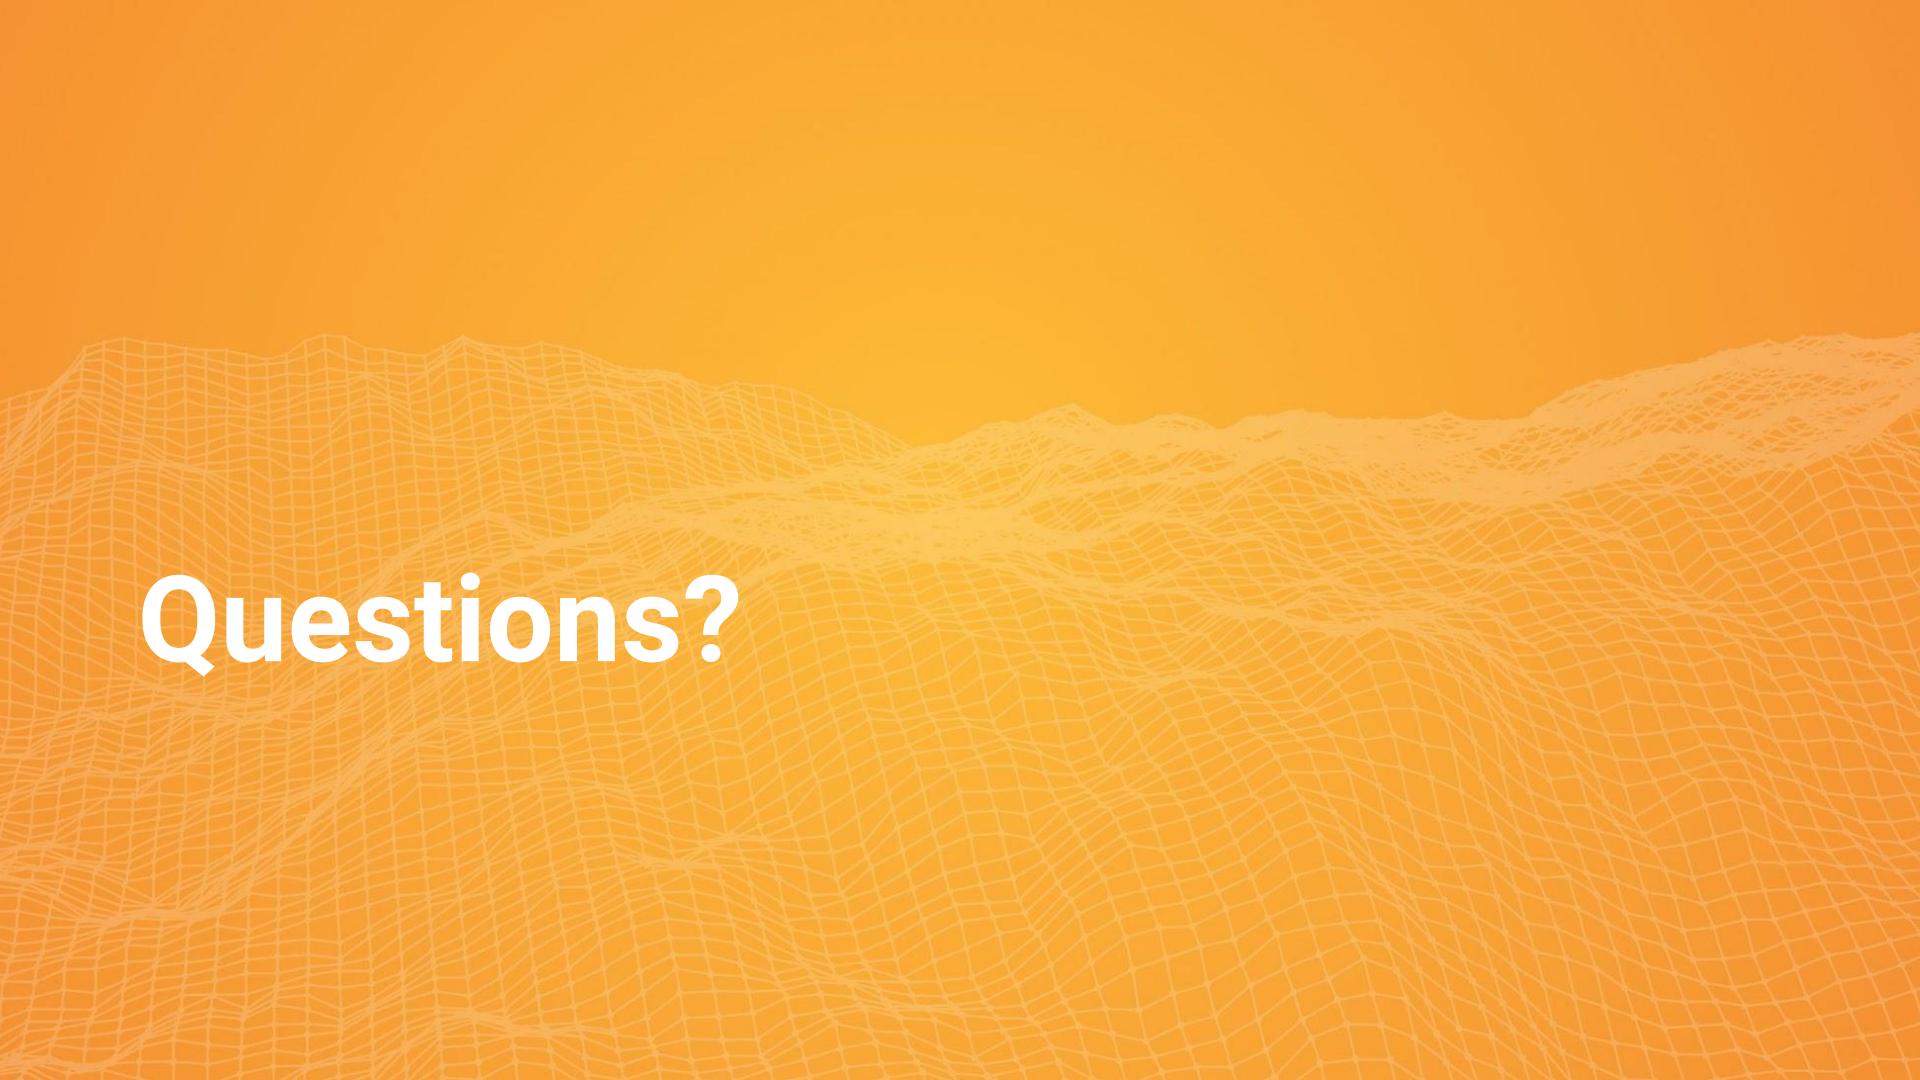
TF Learn Quickstart, goo.gl/uiefRn

TensorFlow for Poets, goo.gl/bVjFIL

ML Recipes, goo.gl/KewA03

TensorFlow and Deep Learning without a PhD, goo.gl/pHeXe7



The background features a stylized, undulating landscape in a warm orange hue. This landscape is composed of a fine, light-colored grid that creates a sense of depth and perspective, resembling a digital or architectural rendering of terrain.

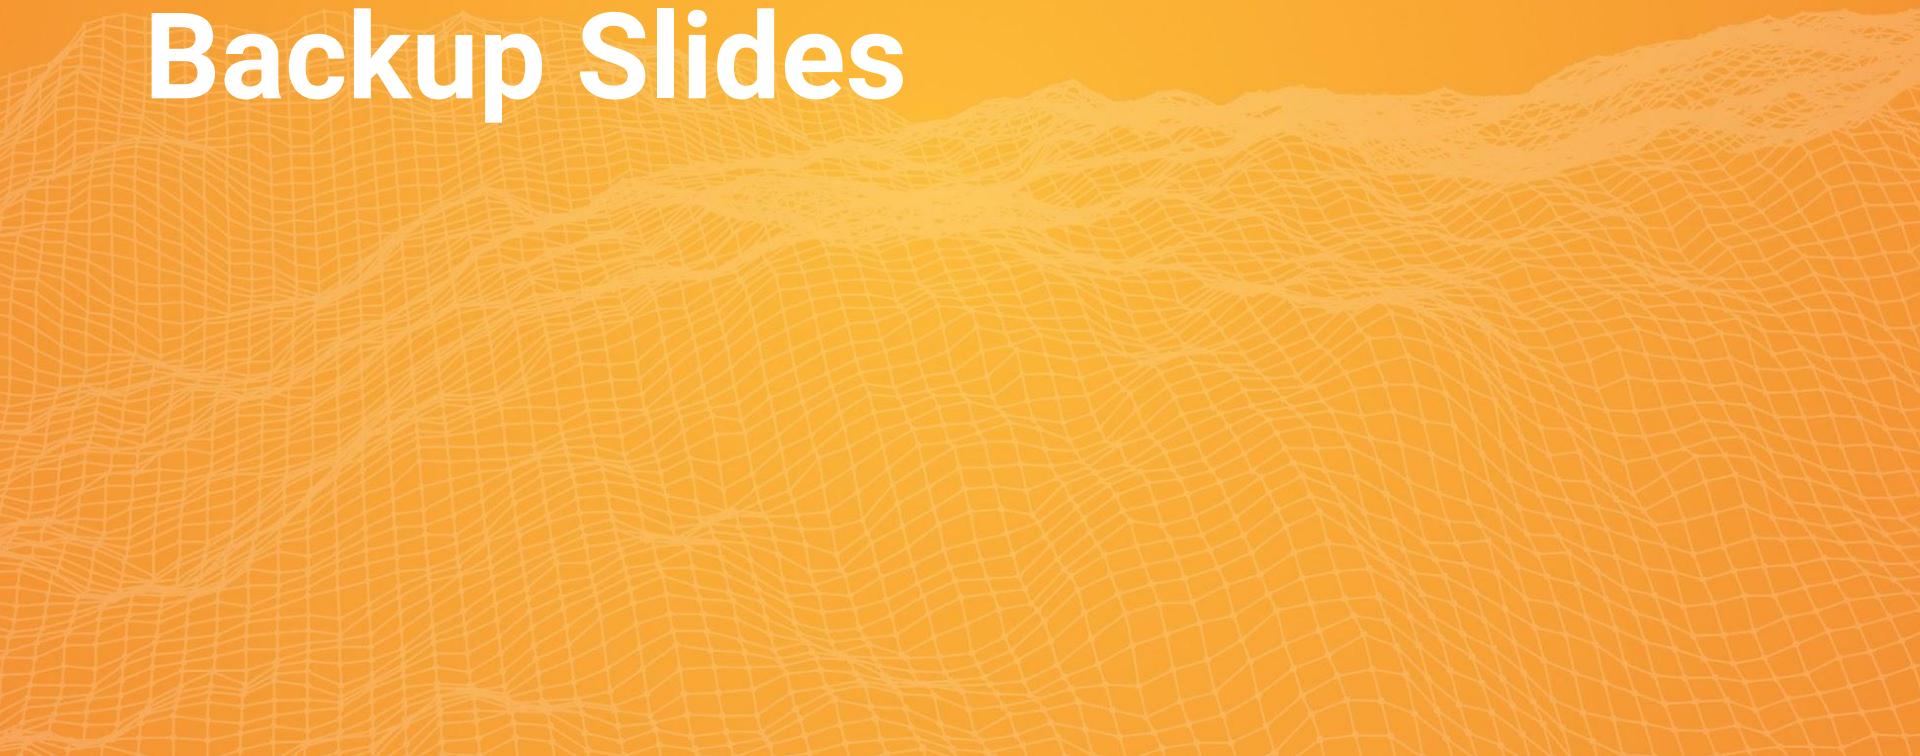
Questions?

Thanks and have fun!



Sherry Q. Moore

Backup Slides





TensorFlow

TensorFlow Playground Demo

Lots of tutorials at tensorflow.org

